

Zero to Hero Study Handbook: NL2SQL Agent

How to use this handbook:

- 1 Read modules in order. Each module builds on the previous one.
- 2 Keep the code open while reading. Every concept here maps to real files/functions in this repo.
- 3 After Module 5, use the checklist to validate end-to-end understanding.

Module 1: Foundations & Architecture

1.1 What this project does

`nl2sql-agent` is a local-first natural-language-to-SQL system. A user asks a plain-English question (CLI or Streamlit UI), the agent generates SQL, validates it for safety, runs it on read-only SQLite or operator-configured PostgreSQL, and returns a natural-language answer plus SQL/result context.

Primary use cases in this repo:

- 1 Local analytics assistant over SQLite (`company.db` by default) or a restricted PostgreSQL schema.
- 2 Safe text-to-SQL experimentation with policy controls (joins, subqueries, aggregates, CTE).
- 3 Multi-provider runtime where Ollama is default and Hugging Face, OpenAI, Anthropic, Gemini, xAI, and Agnes AI are optional.
- 4 Approval-first analysis of session-scoped uploaded SQLite databases.
- 5 Result-based model evaluation across accuracy, safety, latency, retries, and token usage.

1.2 Core paradigms and patterns used

Definitions first, then where they appear in this codebase:

- 1 Modular architecture: Definition: split the system into small modules with one clear job each. In this repo: `config`, `db`, `security`, `llm`, `prompts`, `agent`, `ui`, `utils` under `src/nl2sql_agent/`.
- 2 Factory pattern: Definition: one function returns different implementations behind a shared interface. In this repo: `build_chat_model()` returns provider-specific chat models in `src/nl2sql_agent/llm/factory.py`.
- 3 Stateful workflow (graph/state-machine orchestration): Definition: execution is modeled as nodes and conditional transitions over shared state. In this repo: `NL2SQLAgent.get_workflow()` in `src/nl2sql_agent/agent/workflow.py` using LangGraph.
- 4 Policy-based security (allow-list): Definition: explicitly allow known-safe behavior instead of only blocking known-bad patterns. In this repo: `SQLPolicy + validate_sql()` in `src/nl2sql_agent/security/sql_validator.py`.
- 5 Typed data contracts: Definition: key data objects have explicit types and field names. In this repo: `AgentState` (`agent/state.py`) and `QueryResult` (`db/database.py`).
- 6 Settings singleton with environment binding: Definition: one cached settings object is reused across the app and can be reset in tests. In this repo: `Settings`, `get_settings()`, `reset_settings_cache()` in `config/settings.py`.
- 7 Defense in depth: Definition: multiple independent controls prevent one failure from becoming data loss. In this repo: AST validation, table authorization, preflight compilation, read-only SQLite, progress limits, and redacted audits.

1.3 Architecture description

Key components and interactions:

- 1 Entrypoints:
 - CLI: `src/nl2sql_agent/cli.py` (`main()`, subcommands `ask/config/serve/eval`)
 - UI: `src/nl2sql_agent/ui/streamlit_app.py` (`main()`)
- 2 Model construction:
 - `build_chat_model()` selects provider and returns a `BaseChatModel`.

3 Core agent workflow:

- `NL2SQLAgent` builds a LangGraph with nodes: `fetch_schema` -> `writer` -> `guardian` -> `executor` -> `summarizer`

4 Data access:

- `Database.ensure_schema()` creates/seeds tables.
- `Database.execute()` runs SQL on a hardened read-only connection and returns `QueryResult`.

5 Security gate:

- `prepare_sql()` blocks non-SELECT, multi-statement, dangerous functions, disallowed tables, and policy violations, then enforces the row limit.

6 Prompting:

- SQL-writing and summarization prompts in `prompts/templates.py`.

1.4 Main flow diagram (ASCII)

```

User (CLI ask / Streamlit chat / eval corpus)
  |
  v
build_chat_model(settings, provider?, model?)
  |
  v
NL2SQLAgent(llm, settings)
  |
  v
LangGraph workflow:

START
-> fetch_schema      (db.get_schema_text)
-> writer            (LLM creates sql_query)
-> guardian          (prepare_sql + EXPLAIN QUERY PLAN)
    | safe           | unsafe and retry budget remains
    v               v
  executor          writer
(db.execute)        (LLM SQL rewrite)
    | ok
    | error and retry_count < max_retries
    +-----> writer (retry)

summarizer -> END

Streamlit stops after guardian, shows an editable SQL preview, then calls
execute_prepared(), which validates again before executor.

```

Schema fetching keeps one read-only connection open, scores every table and column identifier once, and loads foreign-key metadata only for the tables that will be sent to the writer.

Module 2: Repository Map

Files a new contributor should learn first (in practical reading order):

File/Directory Path	Primary Responsibility	Key Classes/Functions	Important Configs/Variables
Launch <code>NL2SQL Agent.cmd</code>	Windows double-click startup	checks uv, then runs <code>n12sql-agent serve</code>	locked environment, loopback port 8512, visible logs
<code>pyproject.toml</code>	Packaging, dependencies, scripts, lint/type/test config	<code>[project.scripts]</code> <code>n12sql-agent</code>	<code>requires-python</code> , direct dependencies, transitive constraints, Ruff/ty/Pytest settings
<code>src/n12sql_agent/cli.py</code>	CLI entrypoint for ask/config/serve/eval	<code>main</code> , <code>cmd_ask</code> , <code>cmd_eval</code> , <code>cmd_config</code> , <code>cmd_serve</code>	Provider/model overrides, eval threshold, report and cost options
<code>src/n12sql_agent/ui/streamlit_app.py</code>	Approval-first multipage Streamlit orchestration	<code>main</code> , <code>_resolve_database</code> , <code>_build_agent</code>	Chat context, temporary upload workspace, pending SQL, active database context
<code>src/n12sql_agent/ui/components.py</code>	UI widgets and result rendering	<code>render_sidebar</code> , <code>render_chat_history</code> , <code>render_run_result</code>	Database source, provider/model, results, traces, plans, metrics, CSV download

File/Directory Path	Primary Responsibility	Key Classes/Functions	Important Configs/Variables
src/nl2sql_agent/ui/pages.py	Costs, Sessions, Insights, and Pricing views	render_costs_page, render_sessions_page, render_insights_page, render_pricing_page	Dashboards, saved-session management, rule editing, budgets
src/nl2sql_agent/persistence.py	Versioned local state database	StateStore	Allowlisted messages, approved SQL, usage, price snapshots, plans, metrics; never raw results or secrets
src/nl2sql_agent/ui/database_upload.py	Untrusted upload validation/storage	validate_sqlite_upload, save_sqlite_upload	Extension, size, SQLite header, content digest
src/nl2sql_agent/agent/workflow.py	Full and two-phase workflow routing	NL2SQLAgent, prepare, stream_prepare, execute_prepared, run, stream	Run ID, allowlist, SQL safety, retries, traces, token usage
src/nl2sql_agent/agent/state.py	Typed workflow contract	AgentState	Fields for schema, SQL, retries, raw rows, columns, row_count
src/nl2sql_agent/security/sql_validator.py	SQL AST validation and executable preparation	SQLPolicy, validate_sql, prepare_sql, referenced_tables	Feature toggles, structural limits, table allowlist, enforced LIMIT
src/nl2sql_agent/db/database.py	Read-only SQLite schema/query layer	Database, QueryResult, list_tables, get_schema_text, preflight, execute	Timeout, VM steps, row limit, normalized plans/metrics, schema ranking
src/nl2sql_agent/db/postgres.py	Strict read-only PostgreSQL layer	PostgresDatabase	Restricted role verification, one-schema introspection, timeouts, JSON plans
src/nl2sql_agent/evaluation/runner.py	Result and safety evaluation	EvalCase, EvaluationRunner, EvaluationReport	Reference result comparison, threshold metrics, integrity digest
src/nl2sql_agent/utils/audit.py	Privacy-preserving JSONL audit	AuditLogger, redact_sql, hash_text	Hashed questions, literal-redacted SQL
src/nl2sql_agent/db/seed.py	Demo seed data	SEED_DEPARTMENTS, SEED_EMPLOYEES	Department/employee tuples loaded via INSERT OR IGNORE
src/nl2sql_agent/llm/factory.py	Provider-specific model building and model listing	build_chat_model, list_models, fallback_models, _build_*, _list_ollama, LLMPProviderError	Seven providers; documented hosted reasoning modes; Ollama-only live discovery
src/nl2sql_agent/llm/pricing.py	Per-call hosted-model pricing	PricingRule, UsageRecord, calculate_cost	Effective dates, cache/batch/fast rates, long-context thresholds, exact decimal arithmetic
src/nl2sql_agent/config/settings.py	Central runtime config and model policy	Settings, default_model_for, supported_models_for, validate_model_for, env_var_for	env_prefix="NL2SQL_", provider/model allow-lists, credentials, DB/SQL/log fields
src/nl2sql_agent/prompts/templates.py	Prompt templates and formatting helpers	SQL_WRITER_SYSTEM, SQL_WRITER_USER, SUMMARIZER_SYSTEM, SUMMARIZER_USER, error_section, format_data	Prompt placeholders {schema}, {question}, {error_section}, {sql}, {data}, {error}
src/nl2sql_agent/utils/text.py	SQL/text normalization utilities	strip_sql_fences, truncate	Regex constants _SQL_FENCE_RE, _LEADING_SQL_TOKEN_RE
src/nl2sql_agent/utils/logging.py	Unified Loguru logging setup	configure_logging, get_logger	level, json, stdlib interception for httpx/httpcore/...
tests/unit/test_agent.py	Workflow behavior contract	Routing/retry/fallback tests for node methods and run/stream	Asserts on retry_count, error, final_answer, node names
tests/unit/test_sql_validator.py	Security policy contract	Safety/forbidden/function/subquery/join/CTE/union tests	Confirms allow-list behavior and failure cases

Module 3: Core Execution Flows

This module explains the real operational paths using concrete symbols and data shapes.

3.1 Core data contracts

AgentState `shane (src/nl2sql_agent/agent/state.py)`

```
# TypedDict, total=False
{
    "run_id": str,
    "question": str,
    "schema": str,
    "allowed_tables": list[str],
    "sql_query": str,
    "sql_safe": bool,
    "sql_unsafe_reason": str,
    "result": str,
    "raw_rows": list[tuple[object, ...]],
    "columns": list[str],
    "row_count": int,
    "csv_data": str,
    "truncated": bool,
    "error": str,
    "retry_count": int,
    "max_retries": int,
    "trace": list[dict[str, object]],
    "token_usage": dict[str, int],
    "final_answer": str,
}
```

Important note: `AgentState` is declared with `total=False`. That means each node can return partial updates instead of a full state object.

QueryResult `shane (src/nl2sql_agent/dh/database.py)`

```
QueryResult(
    columns: tuple[str, ...],
    rows: tuple[tuple[object, ...], ...],
    row_count: int,
    truncated: bool
)
```

Sidebar config `navload (src/nl2sql_agent/ui/components.py)`

```
{
    "data_source": str,          # "Demo" | "Upload"
    "provider": Provider,        # ollama | huggingface | openai | anthropic | gemini | xai | agnes
    "model": str | None,
    "api_key": str | None,
}
```

Initial workflow input shape (`run/stream` in `src/nl2sql_agent/agent/workflow.py`)

Both `run()` and `stream()` start with:

```
{
    "run_id": str(uuid4()),
    "question": question,
    "retry_count": 0,
    "max_retries": int(max_retries or self.settings.max_retries),
    "error": "",
    "trace": [],
    "token_usage": {},
    "allowed_tables": sorted(self.allowed_tables),
}
```

3.2 Flow A: CLI ask request

Entrypoint chain:

- 1 `main()` parses args via `build_parser()`.
- 2 `cmd_ask(args)` calls `_build_agent(...)`.
- 3 `_build_agent`:
 - loads `settings = get_settings()`
 - applies CLI overrides (`provider/model/api_key`)
 - builds `llm = build_chat_model(settings)`

- returns `NL2SQLAgent(llm, settings=settings)`
- 4 `cmd_ask` calls `agent.run(args.question)`.
- 5 Prints `result["final_answer"]`; if `--show-sql`, also prints `result["sql_query"]`.

Short code fragment from real path:

```
result = agent.run(args.question)
print(result.get("final_answer", ""))
if args.show_sql:
    print("\n-- SQL ---\n", result.get("sql_query", ""), sep="")
```

Input shape:

- `args.question`: str
- optional overrides: `args.provider`, `args.model`, `args.api_key`

Output shape:

- terminal text; underlying `result` dict includes `final_answer`, `sql_query`, `error`, `columns`, `raw_rows`, `row_count`.

3.3 Flow B: Streamlit chat request

Entrypoint chain:

- 1 `streamlit_app.main()` configures logging, local state, and five native pages.
- 2 `render_sidebar()` selects Demo/Upload/PostgreSQL and the model provider.
- 3 Uploads are validated, stored under a content digest, and opened read-only.
- 4 The user selects allowed tables and whether uploaded sample rows may enter prompts.
- 5 `agent.stream_prepare(user_query)` emits schema, writer, and guardian updates.
- 6 Safe SQL is stored as `pending_query` and shown in an editable text area.
- 7 Run calls `agent.execute_prepared(...)`, which validates the edited SQL again.
- 8 Current result rows and CSV remain in memory. The state store persists only messages, approved SQL, context, usage/cost snapshots, and bounded metrics.
- 9 `render_run_result()` prices each provider-reported model call with the effective local rule and shows plan/runtime warnings.

Short code fragment from real path:

```
for node_name, update in agent.stream_prepare(user_query):
    final_state.update(update)

result = agent.execute_prepared(final_state, sql_query=edited_sql)
```

Input shape:

- `user_query`: str
- sidebar config from `render_sidebar()`

Output shape:

- Editable SQL preview, answer, dataframe, CSV, stage timings, token counts, normalized plan/metrics, warnings, and an estimated hosted-model cost.

Each model call records provider-reported input/output/cache usage. The effective local rule applies its actual request mode and long-context threshold, then freezes a price snapshot with the approved run. Estimates remain local guidance rather than provider invoices.

3.4 Flow C: Agent internals (LangGraph node-by-node)

`NL2SQLAgent.get_workflow()` defines this graph:

- 1 `fetch_schema(state)`:
- ranks allowed tables against the question and includes FK neighbors

- optionally includes bounded sample rows
- returns schema, allowlist, and a timing trace

2 `write_sql(state):`

- reads: `question`, `schema`, `previous error/sql_unsafe_reason`
- builds `SQL_WRITER_USER` prompt
- LLM call with system+user messages
- cleans output via `strip_sql_fences()`
- returns:
- `sql_query`
- incremented `retry_count`
- clears stale `error`, `sql_unsafe_reason`, `result`

3 `check_security(state):`

- runs `prepare_sql(sql, policy, allowed_tables)`
- canonicalizes and enforces the executable LIMIT
- runs `Database.preflight()` using `EXPLAIN QUERY PLAN`
- records a redacted `prepared` or `blocked` audit event

4 `route_after_security(state):`

- if safe -> `"executor"`
- if blocked and retry budget remains -> `"writer"`
- otherwise -> `"summarizer"`

5 `execute_sql(state):`

- calls `self.db.execute(sql)`
- on success returns markdown, rows, columns, count, CSV, truncation state, and timing
- on `sqlite3.Error` returns error string and empty data containers

6 `route_after_execute(state):`

- if `error` and `retry_count < max_retries` -> `"writer"`
- otherwise -> `"summarizer"`

7 `summarize_result(state):`

- builds summarizer prompt from `question`, `sql_query`, `formatted result`, `error`
- LLM generates `final_answer`
- fallback path via `_fallback_answer` if LLM fails or returns empty
- returns `{"final_answer": answer, "result": answer}`

`prepare()/stream_prepare()` use a smaller graph that ends after guardian. `execute_prepared()` re-enters guardian before executor. `run()` and `stream()` continue to use the end-to-end graph for automation compatibility.

3.5 Flow D: Database initialization + seed behavior

Where it happens:

1 For a managed demo in `NL2SQLAgent.__init__`:

- `self.db = Database(...)`
- `self.db.ensure_schema(seed=self.settings.db_seed)`

Injected databases are never initialized or seeded. This is the path used for uploads.

2 `Database.ensure_schema(seed=True):`

- creates parent directories
- executes DDL for `departments` and `employees`
- inserts `SEED_DEPARTMENTS` and `SEED_EMPLOYEES` using `INSERT OR IGNORE`
- guarded by `_init_lock` and `_initialized` flag for idempotency

Tables and seed columns:

1 `departments(dept_id, dept_name, location)`

2 `employees(emp_id, name, salary, dept_id)` with FK `employees.dept_id -> departments.dept_id`

3.6 Flow E: Safety gate behavior details

`prepare_sql()` in `security/sql_validator.py` enforces:

- 1 SQL must parse and be non-empty (`parse_sql`).
- 2 Exactly one statement.
- 3 Top-level must be `SELECT` or set op (`UNION/INTERSECT/EXCEPT`) if policy allows.
- 4 Optional policy bans:
 - subqueries
 - joins
 - aggregates
 - CTEs
 - union
- 5 Dangerous functions are blocked (`load_extension, readfile, writefile, edit, shell, system`).
- 6 Additional banned keyword scan blocks destructive operations.
- 7 JOIN, subquery, and CTE counts stay within configured limits.
- 8 Every referenced physical table belongs to the current allowlist. CTE names are derived sources, but same-name physical tables are still authorized independently through scope traversal.
- 9 Preparation reuses one parsed AST for validation, table discovery, and LIMIT canonicalization; a regression test enforces the single-parse path.
- 10 The returned SQL contains an enforced result LIMIT.

After AST checks, `Database.preflight()` validates identifiers without running the SELECT. Execution uses SQLite URI `mode=ro, query_only`, disabled extensions, `trusted_schema=OFF`, an elapsed-time deadline, and a VM-step cap.

Output contract:

- Success: returns `PreparedSQL(sql=..., tables=...)`.
- Failure: raises `SQLValidationError` with user-facing message.

Why this matters for learners:

- 1 This safety gate is the key trust boundary in the system.
- 2 It runs before query execution.
- 3 It is configurable through `Settings` flags that map to `SQLPolicy`.

Module 4: Setup & Run Guide

This section is static and based on repo files (`README.md`, `pyproject.toml`, `source`).

4.1 Prerequisites

- 1 Python 3.12.10 (from `pyproject.toml`: `requires-python = ">=3.12.10,<3.13"`).
- 2 `uv` for environment/dependency management.
- 3 Optional but default runtime path: local Ollama (`NL2SQL_PROVIDER=ollama`).

4.2 Install on a clean machine

```
git clone <repo-url>
cd natural-language-to-sql-agent
uv venv --python 3.12.10
uv sync --all-groups
uv pip install -e .
```

4.3 Environment configuration (.env or shell env)

Core Settings keys (`src/nl2sql_agent/config/settings.py`)

Use these as `NL2SQL_*` variables:

- 1 `NL2SQL_PROVIDER`
- 2 `NL2SQL_MODEL`
- 3 `NL2SQL_OLLAMA_BASE_URL`
- 4 `NL2SQL_OLLAMA_KEEP_ALIVE`
- 5 `NL2SQL_DB_PATH`
- 6 `NL2SQL_DB_SEED`
- 7 `NL2SQL_DB_MAX_ROWS`
- 8 `NL2SQL_DB_QUERY_TIMEOUT_SECONDS`
- 9 `NL2SQL_DB_MAX_VM_STEPS`
- 10 `NL2SQL_DB_UPLOAD_MAX_MB`
- 11 `NL2SQL_MAX_RETRIES`
- 12 `NL2SQL_LLM_TEMPERATURE`
- 13 `NL2SQL_LLM_MAX_TOKENS`
- 14 `NL2SQL_LLM_REQUEST_TIMEOUT_SECONDS`
- 15 `NL2SQL_SQL_ALLOW_SUBQUERIES`
- 16 `NL2SQL_SQL_ALLOW_JOINS`
- 17 `NL2SQL_SQL_ALLOW_AGGREGATES`
- 18 `NL2SQL_SQL_ALLOW_CTE`
- 19 `NL2SQL_SQL_MAX_JOINS`
- 20 `NL2SQL_SQL_MAX_SUBQUERIES`
- 21 `NL2SQL_SQL_MAX_CTES`
- 22 `NL2SQL_SCHEMA_MAX_TABLES`
- 23 `NL2SQL_AUDIT_ENABLED`
- 24 `NL2SQL_AUDIT_PATH`
- 25 `NL2SQL_LOG_LEVEL`
- 26 `NL2SQL_LOG_JSON`

Provider API key env names used by UI helper

`env_var_for()` in `settings.py` maps:

- 1 `openai -> OPENAI_API_KEY`
- 2 `gemini -> GOOGLE_API_KEY`
- 3 `anthropic -> ANTHROPIC_API_KEY`
- 4 `huggingface -> HF_TOKEN`
- 5 `xai -> XAI_API_KEY`

Minimal `.env` examples:

```
# Local-first default path
NL2SQL_PROVIDER=ollama
NL2SQL_MODEL=phi4-mini:3.8b
NL2SQL_OLLAMA_BASE_URL=http://localhost:11434
NL2SQL_DB_PATH=company.db
NL2SQL_DB_SEED=true

# Cloud provider example (OpenAI)
NL2SQL_PROVIDER=openai
NL2SQL_MODEL=gpt-5.6-luna
OPENAI_API_KEY=your_key_here
NL2SQL_DB_PATH=company.db
```

Hosted requests use medium reasoning. OpenAI allows only `gpt-5.6-luna` and `gpt-5.6-terra`; Anthropic allows `claude-sonnet-5`; Gemini allows `gemini-3.7-flash` and `gemini-3.5-flash-lite`; xAI allows `grok-4.6`. Hugging Face accepts `namespace/model[:routing-policy]` through its direct router and defaults to `openai/gpt-oss-120b:fastest`. Agnes allows only `agnes-2.5-flash` and uses its documented boolean Chat Completions Thinking flag rather than a low/medium/high effort value.

4.4 Typical run commands

Start Streamlit UI

On Windows, double-click `Launch NL2SQL Agent.cmd`. The terminal equivalent is:

```
uv run streamlit run src/nl2sql_agent/ui/streamlit_app.py
```

Ask one question in CLI

```
uv run nl2sql-agent ask "How many employees are in each department?"
uv run nl2sql-agent ask --show-sql "What is the total salary in Engineering?"
```

Show resolved runtime config

```
uv run nl2sql-agent config
```

Serve command wrapper (invokes Streamlit via CLI)

```
uv run nl2sql-agent serve --host 127.0.0.1 --port 8512
```

Run the result and safety corpus

```
uv run nl2sql-agent eval --min-pass-rate 0.8
```

4.5 Database migration/seeding steps

There is no migration framework (no Alembic/Flyway-style migration scripts in this repo).

Current behavior:

- 1 Schema creation is runtime and idempotent via `Database.ensure_schema()`.
- 2 Seeding is controlled by `db_seed` (`NL2SQL_DB_SEED`, default `True`).
- 3 Backward-compatible helper exists: `setup_db(path, seed=True)`.

Module 5: Study Plan & Practice Exercises

5.1 Ordered study plan (zero to hero)

- 1 Read `README.md` sections: architecture, workflow, configuration, CLI/UI.
- 2 Read `src/nl2sql_agent/config/settings.py` to understand runtime knobs.
- 3 Read `src/nl2sql_agent/db/database.py` + `db/seed.py` for actual data model.

- 4 Read `src/nl2sql_agent/security/sql_validator.py` for guardrail logic.
- 5 Read `src/nl2sql_agent/prompts/templates.py` for prompt contracts.
- 6 Read `src/nl2sql_agent/llm/factory.py` for provider abstraction.
- 7 Read `src/nl2sql_agent/agent/state.py` then `agent/workflow.py` linearly.
- 8 Read `src/nl2sql_agent/ui/streamlit_app.py` + `ui/components.py`.
- 9 Read `src/nl2sql_agent/cli.py`.
- 10 Confirm understanding through tests:
 - `tests/unit/test_agent.py`
 - `tests/unit/test_sql_validator.py`
 - `tests/unit/test_database.py`
 - `tests/unit/test_cli.py`

5.2 Practice exercises

- 1 Trace one full run: Starting from `cmd_ask`, list each function called until `final_answer` is printed. File focus: `src/nl2sql_agent/cli.py`, `src/nl2sql_agent/agent/workflow.py`.
- 2 State-contract drill: Write the minimal `AgentState` keys required for `write_sql()` to run without missing context. File focus: `src/nl2sql_agent/agent/state.py`, `src/nl2sql_agent/agent/workflow.py`.
- 3 Security policy scenario: Explain what happens if `allow_joins=False` and the model generates a JOIN query. File focus: `src/nl2sql_agent/security/sql_validator.py`.
- 4 Retry logic reasoning: For `retry_count=2`, `max_retries=3`, and execution error present, what node runs next and why? File focus: `src/nl2sql_agent/agent/workflow.py`.
- 5 Data contract check: What are the exact types for `columns`, `raw_rows`, and `row_count` returned by `execute_sql()`? File focus: `src/nl2sql_agent/db/database.py`, `src/nl2sql_agent/agent/workflow.py`.
- 6 Config mapping task: Map one settings field to an environment variable and describe where it affects runtime behavior. File focus: `src/nl2sql_agent/config/settings.py`.
- 7 UI behavior task: Explain why live `list_models()` discovery is Ollama-only and how hosted model choices are supplied. File focus: `src/nl2sql_agent/ui/streamlit_app.py`, `src/nl2sql_agent/ui/components.py`.
- 8 Extension design: Trace the implemented xAI path from settings and credentials through the factory and sidebar. File focus: `src/nl2sql_agent/config/settings.py`, `src/nl2sql_agent/llm/factory.py`, `src/nl2sql_agent/ui/components.py`.
- 9 Approval flow: Explain why editing SQL in Streamlit cannot bypass the guardian. File focus: `src/nl2sql_agent/ui/streamlit_app.py`, `src/nl2sql_agent/agent/workflow.py`.
- 10 Evaluation scoring: Explain why reference-result comparison accepts multiple correct SQL forms. File focus: `src/nl2sql_agent/evaluation/runner.py`, `src/nl2sql_agent/evaluation/data/demo.jsonl`.

5.3 Solution outlines

- 1 Exercise 1 outline: `main -> parse_args -> cmd_ask -> _build_agent -> get_settings -> build_chat_model -> NL2SQLAgent.run -> get_workflow/invoke -> final dict -> print(final_answer)`.
- 2 Exercise 2 outline: Minimum practical keys are `question`, `schema`; optional prior context is `error` or `sql_unsafe_reason`; `retry_count` defaults to 0 if missing.
- 3 Exercise 3 outline: `check_security` calls `prepare_sql`; the policy raises `SQLValidationError("JOIN clauses are not allowed.")`; the node returns `error + sql_unsafe_reason`; `route_after_security` retries `writer` while budget remains, then ends at `summarizer`.
- 4 Exercise 4 outline: `route_after_execute` checks `err` and `retry < max_retries`; with `2 < 3`, it routes to "writer" for another SQL rewrite attempt.
- 5 Exercise 5 outline: `execute_sql` returns `columns: list[str], raw_rows: list[tuple[object, ...]], row_count: int`.

- 6 Exercise 6 outline: Example: `Settings.db_max_rows` from `NL2SQL_DB_MAX_ROWS` influences `Database(max_rows=...)`; `execute()` fetches `max_rows + 1` then truncates.
- 7 Exercise 7 outline: `list_models()` performs a network call only after `refresh_trigger_ollama` is set. Hosted providers read deterministic choices from the central settings catalog; Hugging Face additionally accepts a custom repository ID.
- 8 Exercise 8 outline: `Provider` includes `xai`; `env_var_for()` maps `XAI_API_KEY`; `_build_xai()` targets the fixed HTTPS endpoint with medium reasoning; the sidebar exposes only `grok-4.6`.

The same trace for Agnes is `Provider` → `AGNES_API_KEY` → `_build_agnes()` → fixed API Hub endpoint with Chat Completions Thinking → `agnes-2.5-flash` in the sidebar.

- 9 Exercise 9 outline: `execute_prepared()` replaces the candidate with the edited SQL, calls `check_security()` again, and only reaches `execute_sql()` when the second validation and preflight succeed.
- 10 Exercise 10 outline: The evaluator executes curated reference SQL read-only and compares result values. SQL aliases, formatting, and equivalent query structure do not have to match; ordered cases still require the requested row order.
- 11 Exercise 11 outline: `write_sql()` and `summarize()` retain per-call provider usage in `AgentState.usage_records`; `calculate_cost()` selects the effective rule and actual mode for each call, including cache and long-context adjustments. Approved runs freeze the price snapshot. Unknown Ollama or custom Hugging Face models remain unpriced until a matching rule is configured.

Verification Checklist

Use this to self-check mastery:

- 1 Can you explain the full node order in `NL2SQLAgent` and both routing decisions?
- 2 Can you describe exactly how unsafe SQL is blocked before execution?
- 3 Can you enumerate the most important `AgentState` keys and what writes each one?
- 4 Can you explain how retries are triggered and when they stop?
- 5 Can you show where schema creation and seed insertion occur?
- 6 Can you explain how provider/model selection flows from UI/CLI into `build_chat_model()`?
- 7 Can you point to where final answer fallback behavior is implemented?
- 8 Can you map at least five environment settings to their runtime effects?
- 9 Can you explain how uploaded databases stay session-scoped and read-only?
- 10 Can you interpret the accuracy, safety, latency, retry, and token fields in an evaluation report?
- 11 Can you explain the UI cost formula and why conditional provider discounts are not applied to aggregate run usage?